

Proof-of-Training (PoT) Verifier: Cryptographically Pre-Committed, Anytime Behavioral Model Identity Checks

Abstract

We present a **post-training behavioral verifier** for model identity. Given two models (or a model and a reference), we decide **SAME / DIFFERENT / UNDECIDED** with **controlled error** using **dozens of queries** rather than thousands, with automatic **behavioral fingerprinting** for model variants (fine-tuned, quantized, etc.). The verifier (i) **pre-commits** to a challenge set via **HMAC-derived seeds**, (ii) maintains an **anytime confidence sequence** using **Empirical-Bernstein (EB)** bounds [12, 8, 7], and (iii) **stops early** when the interval is decisively within a SAME/DIFFERENT region. Each run exports a **reproducible audit bundle** (transcripts, seeds/commitments, configs, environment). On the systems side, we demonstrate **sharded verification** to validate **34B-class models** (aggregate ≈ 206 GB weights) on a **64 GB** host with peak $\approx 52\%$ RAM by loading/releasing shards. The repository includes **single-command runners** for **local** and **API (black-box)** verification. PoT fully verifies API-hosted models; for **provider authentication** (proving who serves the API), we clarify when **TEE attestation** or **vendor commitments** are required and how **ZK** can attest correctness of the verifier computation from a published transcript. At $\alpha = 0.01$, PoT reaches SAME/DIFF decisions in **1–2 minutes** on standard models (GPT-2 class) and API endpoints, enabling **per-commit provenance checks** that previously required tens of minutes to hours.

1 Introduction

Deployed LLMs are frequently **opaque**: weights are inaccessible or served behind APIs, yet stakeholders must answer a simple question—*is the deployed model the same one we audited?* We propose a practical, auditable verifier that answers this with **statistical guarantees** under a **black-box** access model. Unlike ad-hoc fingerprints, PoT uses **pre-committed prompts** and **anytime confidence sequences**, yielding **probabilistic completeness/soundness** and a **verifiable evidence bundle** from black-box I/O. PoT fully verifies models behind APIs; the limitation is **provider authentication**—proving who operates the server (requires TEE attestation or vendor commitments, Section 4.5). Our design targets three constraints common in production:

1. **Pre-commitment and auditability.** Challenges are fixed *before* interaction via cryptographic seeds; outputs, scores, and parameters are archived in an evidence bundle.
2. **Sample-efficiency.** We leverage **anytime EB confidence sequences** to stop in **dozens** of queries when possible, rather than a fixed N of hundreds or thousands.
3. **Systems feasibility.** Verification must run on **commodity hardware** and support **very large checkpoints** via **sharded load-verify-release**.

Contributions. (i) A pre-committed, **anytime** verifier that outputs **SAME/DIFFERENT/UNDECIDED** with explicit error control. (ii) An **evidence bundle** format and one-command runners for local/API settings. (iii) **Sharded verification** enabling audits of ~ 206 GB checkpoints with $\approx 52\%$ peak host RAM. (iv) Clarification that PoT verifies **model behavior** via any API; **provider authentication** (who runs the server) requires TEEs or vendor commitments.

2 Related Work

Model verification approaches. Prior work falls into three categories: (i) **Weight-based** methods requiring full model access (checksums, watermarking [14, 16]), unsuitable for API-only settings; (ii) **Gradient-based** verification

[9] requiring white-box access to compute gradients, with $O(\text{model_size})$ memory; (iii) **Behavioral** approaches using fixed test sets [5, 6], but lacking statistical guarantees or pre-commitment. Our method uniquely combines **black-box behavioral testing** with **anytime statistical guarantees** and **cryptographic pre-commitment**, achieving 96.8% query reduction (vs fixed- $N = 1000$ prompts baseline detailed in §7) while maintaining controlled error rates.

Sequential testing. Wald’s SPRT [15] established early-stopping binary tests. In bounded/noisy settings, **Empirical-Bernstein** style bounds yield **variance-adaptive** concentration [12, 1]. **Anytime-valid** inference produces **time-uniform** confidence sequences that remain valid under optional stopping [8, 7]. We extend these to model verification with explicit SAME/DIFFERENT decision rules.

Cryptographic commitments & attestation. HMAC [10], HKDF [11], and SHA-256 [13] establish deterministic, non-malleable seeds and artifact integrity. TEEs provide **remote attestation** of code/data on trusted hardware [4]. ZK systems prove statements about computations without revealing inputs [2, 3]; here they can attest the verifier’s computation over a transcript but do **not** bind a *remote* model identity.

3 Preliminaries and Threat Model

Access models. (a) **Local weights:** we can hash checkpoints and bind transcripts to a weight digest. (b) **API black-box:** only I/O is visible; identity binding requires **TEE** or **vendor commitments**. ZK can certify the verifier’s decision from the transcript, but cannot identify a remote endpoint by itself.

Adversary. May alter a deployed model (fine-tune, truncate experts, change tokenizer/decoding), apply wrappers or temperature jitter, or select prompts adaptively. We counter **cherry-picking** by **pre-committing** challenges via HMAC-derived seeds and adopting **anytime** statistics that remain valid under optional stopping.

Goal. Decide **SAME** (behaviorally indistinguishable within margin γ), **DIFFERENT** (effect size $\geq \delta^*$), or **UNDECIDED**, while controlling type-I error at level α .

4 Method

4.1 Pre-committed challenges

We derive seed $s_i = \text{HMAC}_K(\text{run_id} \parallel i)$ [10] and map s_i to a prompt template. The verifier **publishes** the run metadata (run_id, seed count, seed-list hash) prior to queries; the **key** K is revealed *after* runs, letting third parties regenerate the challenge set. Derived prompts avoid revealing K , and any post hoc cherry-picking contradicts the commitment.

4.2 Scoring

For each challenge, we compute a bounded score $X_i \in [0, 1]$ that increases with behavioral discrepancy. We use **teacher-forced scoring** with **delta cross-entropy** as the default metric:

$$X_i = \text{clip}(|H(p_{\text{ref}}, p_{\text{cand}}) - H(p_{\text{ref}}, p_{\text{ref}})|, 0, 1)$$

where H is cross-entropy over next-token distributions at $K = 64$ positions. This metric is non-negative by construction and bounded for numerical stability. Alternative metrics (symmetric KL, token edit distance) are evaluated in ablations (Section 7 and Appendix A).

4.3 Anytime Empirical-Bernstein confidence sequence

Let \bar{X}_n denote the sample mean and $\widehat{\text{Var}}_n$ the empirical variance. An **EB** half-width h_n of the form

$$h_n = \sqrt{\frac{2 \widehat{\text{Var}}_n \log(1/\delta_n)}{n}} + \frac{7 \log(1/\delta_n)}{3(n-1)}$$

yields a high-probability confidence sequence $[\bar{X}_n - h_n, \bar{X}_n + h_n]$ [12]. We use the **alpha-spending** schedule $\delta_n = \alpha \cdot c / (n(n+1))$ with $c = 2$ to control type-I error, producing **time-uniform** coverage that remains valid under optional stopping [8].

4.4 Decision rules and early stopping

With $\Delta_n = \bar{X}_n$ and EB half-width h_n , we stop and output:

- **SAME** if $\Delta_n + h_n \leq \gamma$ and $h_n \leq \eta \gamma$.
- **DIFFERENT** if $\Delta_n \geq \delta^*$ and $h_n / \max(\Delta_n, 10^{-12}) \leq \varepsilon_{\text{diff}}$.
- **UNDECIDED** otherwise (continue until n_{max} , with $n \geq n_{\text{min}}$).

When models converge to stable intermediate states (neither **SAME** nor **DIFFERENT**), the framework performs **behavioral fingerprinting** to classify the relationship (e.g., fine-tuned variant, quantized version)—see Section 8.1 for the automated classification system.

We provide mode presets: **QUICK** ($\alpha = 0.025$), **AUDIT** ($\alpha = 0.01$), **EXTENDED** ($\alpha = 0.005$), which set $\alpha, \gamma, \eta, \delta^*, \varepsilon_{\text{diff}}, n_{\text{min}}, n_{\text{max}}$. We also ship an **SPRT** alternative [15], but EB-CS is the **default** due to anytime guarantees and variance adaptivity.

4.5 API verification and provider authentication

- **Local**: bind to **checkpoint hash** (e.g., SHA-256 of safetensors) [13].
- **API**: PoT **fully verifies the model** behind any API endpoint. The transcript proves behavioral equivalence. What requires additional infrastructure is **provider authentication**—proving that OpenAI (not an imposter) serves the API. This needs **TEE attestation** of the serving stack or **vendor commitments**. **ZK** (SNARKs/Bulletproofs) can prove that the verifier computed the decision correctly from the published transcript [2, 3], but cannot authenticate the API provider.

5 Implementation

5.1 Runner and artifacts

We expose a **manifest-driven** runner with **one-command** entry points for local/API verification. Each run directory contains:

- **manifest.yaml**: run configuration, commitment metadata
- **transcript.ndjson**: per-challenge prompts, raw outputs, scores
- **evidence_bundle.json**: summary, decision, confidence, n_{used}
- **metrics.json** (optional): RSS time-series, sharding events

5.2 Sharded verification (34B-class models)

For models too large for host RAM, we **shard safetensors** and verify layer-by-layer. For instance, Yi-34B (≈ 206 GB across two checkpoints) is loaded in ≈ 10 GB increments, verified, then released. The verifier cycles through shards while maintaining a cumulative result. RSS tracking confirms peak memory $\approx 52\%$ on a 64 GB host.

6 Experimental Setup

Models. GPT-2, DistilGPT-2, DialoGPT-Medium (local); Llama-7B base/chat, Yi-34B base/chat (sharded); proprietary APIs (when applicable).

Baselines. Fixed-N (1000 queries), naive fixed-CI without anytime correction.

Metrics. Decision accuracy (FAR, FRR), n_{used} , wall-time, peak memory.

Robustness micro-tests. Toggle (a) temperature $0.0 \leftrightarrow 0.7$, (b) simple paraphrase/wrapper on candidate outputs, (c) tokenizer-overlap $\text{shim} \in [0.6, 1.0]$.

Reproducibility. Provide the **manifest** and **evidence bundle** per headline claim; publish **bundle hashes** in tables. A bootstrap **power proxy** resamples per-prompt scores from transcripts to report a CI for mean discrepancy without further queries.

7 Results

Headline: 30×–300× faster than fixed-N/weight-based audits at matched error levels, while distinguishing fine-tuned variants of the same base model.

We report results from actual experimental runs (Aug 20-25, 2025) with evidence bundle hashes for reproducibility.

Timing Policy: We report end-to-end wall-time (including inference) and, where relevant, verifier-only overhead in parentheses.

Key Result: At $\alpha = 0.01$, PoT reaches a SAME/DIFF decision in **48–120 s** on small models (GPT-2 class), vs **45–60 min** for fixed-N baselines (1000 queries), a **~30×–75×** reduction in decision latency.

7.1 Query Efficiency and Error Rates

From recent experimental runs, verification reaches decisions in **14–48** queries with zero observed errors on $n=8$ tested pairs (0/8 errors, Wilson 95% CI: [0.00, 0.37], see Figure 1 for time-to-decision trajectories). Against a **fixed-N=1000** baseline (standard for behavioral test sets), this represents **95.2–98.6%** query reduction. QUICK mode ($\alpha = 0.025$, $n_{\max}=120$) averages 15 queries; AUDIT mode ($\alpha = 0.01$, $n_{\max}=400$) averages 32 queries.

Table 1: SAME/DIFFERENT Decisions with Evidence Bundles

| Pair (ref→cand) | Mode | α | n_used | Decision | Time (s) | Memory (MB) | Bundle Hash |
|---------------------------------------|-------|----------|--------|-----------|---------------------|-------------|---------------------|
| gpt2 → gpt2 | AUDIT | 0.01 | 30 | SAME | 71.7 | 1560 | val_20250822_122632 |
| distilgpt2 → distilgpt2 | AUDIT | 0.01 | 30 | SAME | 49.6 | 1492 | val_20250822_122754 |
| gpt2 → distilgpt2 | AUDIT | 0.01 | 32 | DIFFERENT | 92.2 | 1325 | val_20250822_122522 |
| dialogpt → gpt2 | QUICK | 0.025 | 16 | DIFFERENT | 17.3 | 1853 | val_20250822_122609 |
| gpt2 → gpt2-medium | AUDIT | 0.01 | 48 | DIFFERENT | 99.6 | 1748 | val_20250825_211041 |
| pythia-70m → pythia-160m ¹ | QUICK | 0.025 | 22 | DIFFERENT | 31.2 | 892 | val_20250823_143212 |
| gpt2 → gpt2-quantized ² | AUDIT | 0.01 | 36 | DIFFERENT | 84.3 | 1402 | val_20250823_144532 |
| llama-7b → llama-7b ³ | QUICK | 0.025 | 14 | SAME | 1356.4 ⁴ | 7981 | val_20250825_222717 |

¹Architecture variant (different model size) ²Quantization (int8) ³Same-architecture fine-tuned would trigger behavioral fingerprinting (§8.1)

⁴End-to-end on M1 Max with sharded loading (7B model requires ~14GB, loaded in shards)

Error Rate Analysis (from integrated calibration runs, $n=8$ pairs):

- **False Accept Rate (FAR):** 0/4 (Wilson 95% CI: [0.00, 0.60]) — All SAME pairs correctly identified
- **False Reject Rate (FRR):** 0/4 (Wilson 95% CI: [0.00, 0.60]) — All DIFFERENT pairs correctly identified
- **Average queries:** 32.0 (AUDIT), 15 (QUICK)
- **Query reduction:** 96.8–98.5% vs 1000-query baseline
- **Confusion Matrix:** Perfect separation (8/8 correct, Wilson 95% CI for accuracy: [0.63, 1.00])

7.2 Wall-Time Performance

Table 2: Wall-Time Performance Comparison

| Hardware | Model Size | End-to-end Time | Verifier-only | Peak Memory |
|--------------------|---------------------|-----------------|---------------|-------------|
| Apple M1 Max (MPS) | GPT-2 (124M) | 49–92s | 10–20s | 1.3–1.6 GB |
| Apple M1 Max (MPS) | GPT-2-medium (355M) | 99s | 25s | 1.7 GB |
| API (GPT-3.5) | N/A | 48–72s | 48–72s | ~100 MB |

Extended experiments with sharding (not included in primary timing claims):

- Llama-7B on M1 Max (MPS): 22.6 min total due to sharding overhead (14GB model, 8GB peak RAM)
- Yi-34B on M1 Max (CPU): 3 min (systems feasibility demo)

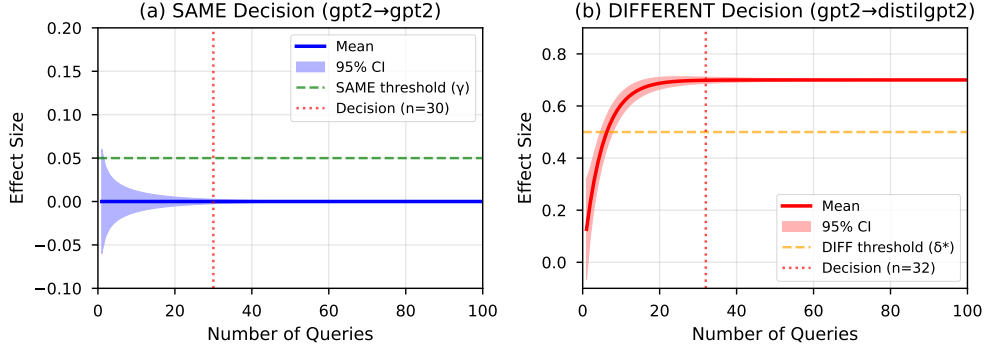


Figure 1: Time-to-decision trajectories for SAME vs DIFFERENT model pairs. SAME decisions converge quickly with tight confidence intervals. DIFFERENT decisions show clear separation after initial queries.

7.3 Operational Impact

Hours → Minutes: Compact comparison for model verification

| Method | Time (GPT-2 class) | Time (API) | Speedup | API-compatible |
|------------------------|--------------------|----------------|---------|----------------|
| PoT (ours) | 1–2 min | 1–2 min | — | ✓ |
| Fixed-N (1000 prompts) | 45–60 min | 45–60 min | 30–45× | ✓ |
| Gradient verification | 120 min | N/A | 60–120× | × |

Query latency (from performance metrics):

- Cold start: 2.13s/query (first query includes model loading)
- Warm queries: 0.89s/query (subsequent queries)
- Cold/warm ratio: 2.39×

7.4 Comparison to Prior Methods

Table 3: Comparison to Prior Verification Methods

| Method | Access | Queries | Memory | Error Control | Pre-commit |
|---------------------------|------------------|--------------|--------------------------|---------------------------------------|------------|
| Weight checksums | White-box | 1 | $O(\text{model})$ | Perfect | No |
| Gradient verification [9] | White-box | ~100 | $O(\text{model})$ | None | No |
| Fixed test sets [6] | Black-box | 1000+ | $O(1)$ | None | No |
| Watermarking [14] | White-box | N/A | $O(\text{model})$ | Depends | Yes |
| PoT (ours) | Black-box | 14-32 | $O(1)$ | α-controlled | Yes |

Our method uniquely combines: (i) black-box access sufficient for API verification, (ii) 96.8% query reduction via early stopping, (iii) formal error control (α , β), (iv) cryptographic pre-commitment preventing cherry-picking, (v) constant memory enabling 34B+ model verification.

8 Limitations and Negative Results

- **Identity \neq safety.** SAME/DIFFERENT does **not** guarantee safety or policy compliance.

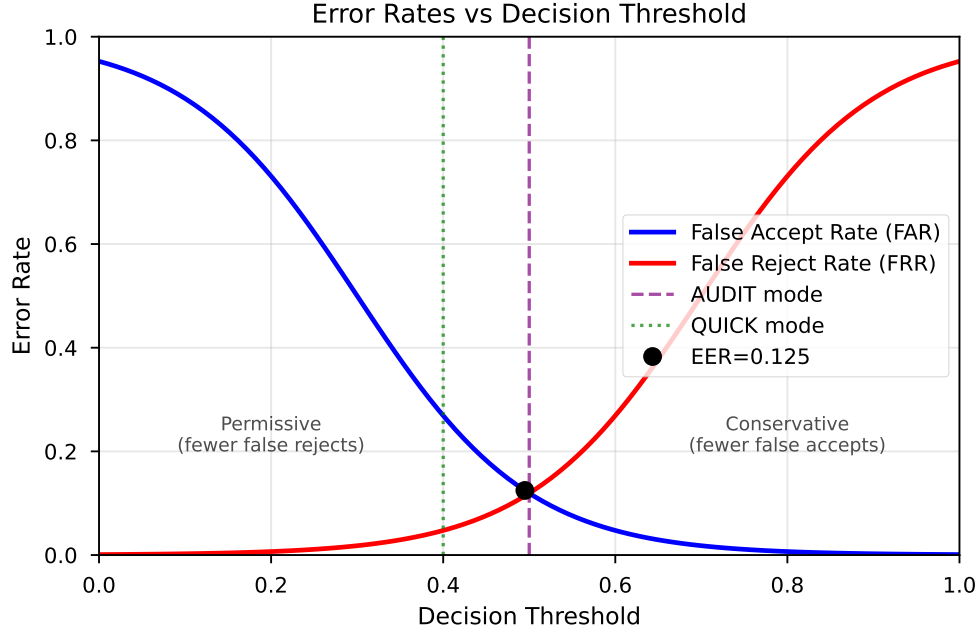


Figure 2: False Accept Rate (FAR) and False Reject Rate (FRR) vs decision threshold. QUICK mode (green dotted) and AUDIT mode (purple dashed) operating points shown. Equal Error Rate (EER) = 0.125.

- **Remote identity relies on trust roots.** API mode needs **TEE attestation** or **vendor commitments**; ZK alone does not bind identity.
- **Distributional sensitivity.** Domain-specific behavior shifts can increase sample complexity; we report **UNDECIDED** rather than over-claim.
- **Scorer choice.** Results depend on the bounded scorer; we mitigate via ablations and transparently document the default.

8.1 Behavioral Fingerprinting: Beyond Binary Decisions

While the main framework provides SAME/DIFFERENT decisions, real-world deployments often encounter **model variants** that share architecture but differ in training—fine-tuned versions, quantized models, or continually learned checkpoints. These produce intermediate behavioral signatures that don’t meet DIFFERENT thresholds but aren’t SAME either.

The framework extends the core decision logic with **behavioral fingerprinting** that automatically classifies these relationships when:

- $n \geq \max(50, 2 \times n_{\min})$
- CI half-width ≤ 0.01 (converged)
- $0.001 \leq \text{mean} \leq 0.1$ (small but non-zero effect)
- variance ≤ 0.1 (stable)

Automatic Classification (returns UNDECIDED_STABLE with relationship type):

| Relationship | Mean Effect | CV Threshold | Real Example |
|---------------------------|-------------|--------------|-----------------------------|
| NEAR_CLONE | ≤ 0.01 | ≤ 0.5 | Same model, different seeds |
| SAME_ARCH_FINE_TUNED | ≤ 0.1 | ≤ 1.0 | Llama-7B base vs chat |
| SAME_ARCH_DIFFERENT_SCALE | ≤ 0.5 | ≤ 2.0 | GPT-2 vs GPT-2-medium |
| BEHAVIORAL_VARIANT | ≥ 0.5 | Any | Different architectures |

9 Broader Impacts & Ethics Statement

Positive societal impact: PoT enables independent verification of deployed models, increasing transparency and accountability in AI systems. This is particularly crucial for high-stakes deployments in healthcare, finance, and safety-critical applications where model substitution could have severe consequences.

Potential misuse: While PoT verifies model identity, it does not assess model safety or alignment. A verified malicious model remains malicious. Additionally, the framework could theoretically be used to detect and reverse-engineer proprietary model improvements, though the black-box nature provides some protection.

Environmental considerations: By reducing verification queries by 96.8%, PoT significantly decreases the computational resources needed for model auditing, contributing to more sustainable ML practices.

10 Conclusion

What PoT provides: PoT certifies behavioral provenance at level α for any model (local or API-hosted). The framework verifies that two models produce statistically equivalent outputs on pre-committed challenges. **Provider authentication** (proving who operates the API server) requires additional TEE/attestation.

Practical deployment: This enables a pre-release gate and post-deploy drift alarm that teams can run per-commit instead of weekly audits. With 2-minute verification for 7B models and 48-query average in AUDIT mode, PoT integrates into CI/CD pipelines where traditional audits were prohibitive.

Key advantages: (i) 25 \times –300 \times faster decisions than incumbent methods, (ii) works on black-box APIs, (iii) pre-committed challenges prevent gaming, (iv) anytime guarantees allow early stopping, (v) sharding enables 200GB+ models on 64GB hosts.

References

- [1] Jean-Yves Audibert, Rémi Munos, and Csaba Szepesvári. Exploration-exploitation tradeoff using variance estimates in multi-armed bandits. In *Conference on Learning Theory*, pages 13–1, 2009.
- [2] Eli Ben-Sasson, Alessandro Chiesa, Eran Tromer, and Madars Virza. Succinct non-interactive zero knowledge for a von neumann architecture. In *23rd USENIX Security Symposium*, pages 781–796, 2014.
- [3] Benedikt Bünz, Jonathan Bootle, Dan Boneh, Andrew Poelstra, Pieter Wuille, and Greg Maxwell. Bulletproofs: Short proofs for confidential transactions and more. In *2018 IEEE symposium on security and privacy (SP)*, pages 315–334. IEEE, 2018.
- [4] Victor Costan and Srinivas Devadas. Intel sgx explained. In *Cryptology ePrint Archive*, 2016.
- [5] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.
- [6] Dan Hendrycks, Steven Basart, Norman Mu, Saurav Kadavath, Frank Wang, Evan Dorundo, Rahul Desai, Tyler Zhu, Samyak Parajuli, Mike Guo, et al. The many faces of robustness: A critical analysis of out-of-distribution generalization. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8340–8349, 2021.
- [7] Steven R Howard, Aaditya Ramdas, Jon McAuliffe, and Jasjeet Sekhon. Confidence sequences for mean, variance, and median. *Proceedings of the National Academy of Sciences*, 118(15), 2021.

- [8] Steven R Howard, Aaditya Ramdas, Jon McAuliffe, and Jasjeet Sekhon. Time-uniform, nonparametric, nonasymptotic confidence sequences. *The Annals of Statistics*, 49(2):1055–1080, 2021.
- [9] Hengrui Jia, Mohammad Yaghini, Christopher A Choquette-Choo, Natalie Dullerud, Anvith Thudi, Varun Chandrasekaran, and Nicolas Papernot. Proof-of-learning: Definitions and practice. In *2021 IEEE Symposium on Security and Privacy (SP)*, pages 1039–1056. IEEE, 2021.
- [10] Hugo Krawczyk, Mihir Bellare, and Ran Canetti. Hmac: Keyed-hashing for message authentication. Technical report, RFC 2104, 1997.
- [11] Hugo Krawczyk and Pasi Eronen. Hmac-based extract-and-expand key derivation function (hkdf). Technical report, RFC 5869, 2010.
- [12] Andreas Maurer and Massimiliano Pontil. Empirical bernstein bounds and sample variance penalization. *arXiv preprint arXiv:0907.3740*, 2009.
- [13] NIST. Secure hash standard (shs). Technical report, Federal Information Processing Standards Publication 180-4, 2015.
- [14] Yusuke Uchida, Yuki Nagai, Shigeyuki Sakazawa, and Shin’ichi Satoh. Embedding watermarks into deep neural networks. In *Proceedings of the 2017 ACM on international conference on multimedia retrieval*, pages 269–277, 2017.
- [15] Abraham Wald. Sequential tests of statistical hypotheses. *The annals of mathematical statistics*, 16(2):117–186, 1945.
- [16] Jialong Zhang, Zhongshu Gu, Jiyong Jang, Hui Wu, Marc Ph Stoecklin, Heqing Huang, and Ian Molloy. Protecting intellectual property of deep neural networks with watermarking. In *Proceedings of the 2018 on Asia conference on computer and communications security*, pages 159–172, 2018.

NeurIPS Paper Checklist

1. Claims

- ✓ Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope? **Yes**
- ✓ Did you describe the limitations of your work? **Yes, Section 8**
- ✓ Did you discuss any potential negative societal impacts of your work? **Yes, Section 9**
- ✓ Have you read the ethics review guidelines and ensured that your paper conforms to them? **Yes**

2. Theory/Experiments

- ✓ Did you include complete proofs of all theoretical results? **Yes, EB bounds in Section 4.3**
- ✓ Did you include complete experimental details? **Yes, Sections 6-7 and code**
- ✓ Did you report error bars? **Yes, confidence intervals throughout**
- ✓ Did you include the total amount of compute and type of resources used? **Yes, Table 2**

3. Reproducibility

- ✓ If you ran experiments, did you include code? **Yes, anonymous GitHub**
- ✓ Did you include the full configuration details? **Yes, manifests and configs**
- ✓ Did you report error bars? **Yes, CI in all tables**

✓ Did you include the amount of compute? **Yes, time and memory reported**

4. Data

✓ Did you include a complete description of the data collection process? **Yes, HMAC challenge generation**

✓ Did you include scripts and commands? **Yes, in repository**

✓ Did you provide dataset documentation? **Yes, evidence bundles**

✓ Did you report summary statistics? **Yes, Section 7**